

XGBoost Model for Predicting Staff Termination

Author: Omkar Masurekar

Description: This notebook demonstrates the process of feature engineering, model training, and prediction using the XGBoost algorithm to identify the likelihood of staff termination. The workflow includes data preprocessing, handling categorical variables, and evaluating model performance to provide actionable insights into employee attrition risk.

```
In [392]: 1 import pandas as pd
          2 import numpy as np
          3 pd.set_option('display.max_columns', None)

In [393]: 1 df= pd.read_csv(r"./cleaned_test_csv.csv")

In [394]: 1 df
```

Out[394]:

Unnamed: 0	FirstName	LastName	StartDate	ExitDate	PositionTitle	Supervisor	
0	Aaden	Mercer	2023-07-26	NaN	Senior Research Fellow	Victoria Jacobs	aaden
1	Aaliyah	Watts	2020-05-09	NaN	Senior Lecturer	Jared Whitehead	aaliya
2	Aarav	Espinoza	2019-11-11	NaN	Officer	David Ali	aarav.e
3	Aaron	Tapia	2023-01-20	NaN	Officer	Michael Reed	aaron
4	Aaron	Weber	2020-04-06	2023-07-24	Lecturer	Melanie Garcia	aaron
...
2995	Zoey	Page	2019-05-06	2022-04-09	Professor	William Newman	zoey
2996	Zoey	Spence	2020-05-06	NaN	Senior Lecturer	Ian Price	zoey
2997	Zoie	Logan	2019-04-19	NaN	Lecturer	Heidi Terry	zoie
2998	Zoie	Mercado	2020-12-01	NaN	Lecturer	David Young	zoie.n
2999	Zoie	Rowland	2022-02-08	2022-07-22	Senior Lecturer	Mrs. Rachel Henry	zoie

3000 rows × 42 columns

```
In [395]: 1 df.columns
```

```
Out[395]: Index(['Unnamed: 0', 'FirstName', 'LastName', 'StartDate', 'ExitDate',  
                'PositionTitle', 'Supervisor', 'ADEmail', 'Faculty', 'EmployeeStatus',  
                'EmployeeType', 'TerminationType', 'DOB', 'JobFunction', 'GenderCode',  
                'Current Employee Rating', 'Employee ID', 'Survey Date',  
                'Engagement Score', 'Satisfaction Score', 'Work-Life Balance Score',  
                'StaffType', 'EmploymentDuration_months', 'TerminationYear',  
                'TerminationMonth', 'TerminationYearMonth', 'TenureGroup',  
                'EmploymentStatus', 'StatusValidated', 'TerminationCategory',  
                'VoluntaryTermination', 'InvoluntaryTermination', 'ActiveStaff',  
                'RetirementCategory', 'TerminatedStaff', 'RetainedOver2Years',  
                'Total_Survey_Score', 'PerformanceGroup', 'EngagementGroup',  
                'SatisfactionGroup', 'Work-Life_BalanceGroup', 'Total_Survey_Group'],  
              dtype='object')
```

```
In [396]: 1 df['DOB'] = pd.to_datetime(df['DOB'], dayfirst=True, errors='coerce')  
2 df['StartDate'] = pd.to_datetime(df['StartDate'], errors='coerce')  
3 df['Age'] = (pd.Timestamp.today() - df['DOB']).dt.days // 365  
4 df['TerminationMonth'] = df['TerminationMonth'].fillna('Active')
```

```
In [397]: 1 df['TenureGroup'].unique()
```

```
Out[397]: array(['1-3yrs', '5-10yrs', '3-5yrs', '<1yr', '0'], dtype=object)
```

Since JobFunction and PositionTitle have a large number of categorical values, we applied label encoding to convert each category into an integer value. This approach helps simplify these features for the model and improves training efficiency.

```
In [398]: 1 from sklearn.preprocessing import LabelEncoder  
2 import pandas as pd  
3 import numpy as np  
4 from xgboost import XGBClassifier  
5 from sklearn.model_selection import train_test_split  
6 from sklearn.metrics import classification_report, accuracy_score  
7 from sklearn.preprocessing import OneHotEncoder  
8 from sklearn.compose import ColumnTransformer  
9 from sklearn.pipeline import Pipeline  
10  
11  
12 for col in ['JobFunction', 'PositionTitle']:  
13     le = LabelEncoder()  
14     df[col] = df[col].fillna('Unknown') # Handle missing values  
15     df[col + '_LE'] = le.fit_transform(df[col])
```

```
In [399]: 1 tenure_order = ['0', '<1yr', '1-3yrs', '3-5yrs', '5-10yrs', '10+yrs']
2
3 from sklearn.preprocessing import OrdinalEncoder
4
5 enc = OrdinalEncoder(categories=[tenure_order], dtype=int)
6 df['TenureGroup_'] = enc.fit_transform(df[['TenureGroup']])
7
```


```
In [400]: 1 df.groupby(['TenureGroup', 'TenureGroup_']).size()
```

```
Out[400]: TenureGroup  TenureGroup_
0                0                6
1-3yrs          2            1020
3-5yrs          3             725
5-10yrs         4             538
<1yr           1             711
dtype: int64
```

```
In [401]: 1 df.columns
```

```
Out[401]: Index(['Unnamed: 0', 'FirstName', 'LastName', 'StartDate', 'ExitDate',
                'PositionTitle', 'Supervisor', 'ADEmail', 'Faculty', 'EmployeeStatus',
                'EmployeeType', 'TerminationType', 'DOB', 'JobFunction', 'GenderCode',
                'Current Employee Rating', 'Employee ID', 'Survey Date',
                'Engagement Score', 'Satisfaction Score', 'Work-Life Balance Score',
                'StaffType', 'EmploymentDuration_months', 'TerminationYear',
                'TerminationMonth', 'TerminationYearMonth', 'TenureGroup',
                'EmploymentStatus', 'StatusValidated', 'TerminationCategory',
                'VoluntaryTermination', 'InvoluntaryTermination', 'ActiveStaff',
                'RetirementCategory', 'TerminatedStaff', 'RetainedOver2Years',
                'Total_Survey_Score', 'PerformanceGroup', 'EngagementGroup',
                'SatisfactionGroup', 'Work-Life_BalanceGroup', 'Total_Survey_Group',
                'Age', 'JobFunction_LE', 'PositionTitle_LE', 'TenureGroup_'],
                dtype='object')
```

```
In [402]: 1 features = ['PositionTitle_LE', 'Faculty', 'EmployeeType', 'JobFunction',
2                 'StaffType', 'EmploymentDuration_months', 'Engagement Score', 'S
3
```



```
In [403]: 1 features = list(dict.fromkeys(features))
2 categorical_features = [col for col in features if df[col].dtype == 'object']
3 numerical_features = [col for col in features if col not in categorical_features]
4
```

```
In [404]: 1 numerical_features
```

```
Out[404]: ['PositionTitle_LE',  
           'JobFunction_LE',  
           'Current Employee Rating',  
           'EmploymentDuration_months',  
           'Engagement Score',  
           'Satisfaction Score',  
           'Work-Life Balance Score',  
           'TenureGroup_',  
           'Total_Survey_Score',  
           'Age']
```

```
In [405]: 1 categorical_features
```

```
Out[405]: ['Faculty', 'EmployeeType', 'GenderCode', 'StaffType']
```

```
In [406]: 1 preprocessor = ColumnTransformer(  
2     transformers=[  
3         ('cat', OneHotEncoder(handle_unknown='ignore', sparse=False), c  
4         ('num', 'passthrough', numerical_features)  
5     ]  
6 )  
7 pipeline = Pipeline(steps=[  
8     ('preprocessor', preprocessor),  
9     ('classifier', XGBClassifier(use_label_encoder=False, eval_metric='  
10 ])
```

```
In [407]: 1 X_train, X_test, y_train, y_test = train_test_split(  
2     df[features], df['TerminatedStaff'], test_size=0.2, random_state=42  
3
```

```
In [408]: 1 import joblib  
2 pipeline.fit(X_train, y_train)  
3 joblib.dump(pipeline, './xgb_pipeline_model.pkl')  
4 #pipeline = joblib.load('xgb_pipeline_model.pkl')  
5 y_pred = pipeline.predict(X_test)
```

C:\Users\omas0005\AppData\Local\anaconda3\Lib\site-packages\sklearn\preprocessing_encoders.py:868: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

warnings.warn(

C:\Users\omas0005\AppData\Local\anaconda3\Lib\site-packages\xgboost\training.py:183: UserWarning: [10:56:48] WARNING: C:\actions-runner_work\xgboost\xgboost\src\learner.cc:738:

Parameters: { "use_label_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)

In [409]: 1 `print(classification_report(y_test, y_pred))`

	precision	recall	f1-score	support
0	0.82	0.86	0.84	293
1	0.86	0.82	0.84	307
accuracy			0.84	600
macro avg	0.84	0.84	0.84	600
weighted avg	0.84	0.84	0.84	600

In [410]: 1 `# Predict the probability`

In [411]: 1 `df_active = df[df['TerminatedStaff'] == 0]`

In [412]:

```
1 X_active = df_active[features]
2 pred_leave = pipeline.predict(X_active)
3 prob_leave = pipeline.predict_proba(X_active)[:, 1]
4 df_active = df_active.copy()
5 df_active['Predicted_Terminated'] = pred_leave
6 df_active['Probability_Terminated'] = prob_leave
```

In [413]:

```
1 df_active
```

Out[413]:

Unnamed: 0		FirstName	LastName	StartDate	ExitDate	PositionTitle	Supervisor	
0	0	Aaden	Mercer	2023-07-26	NaN	Senior Research Fellow	Victoria Jacobs	aaden
1	1	Aaliyah	Watts	2020-05-09	NaN	Senior Lecturer	Jared Whitehead	aaliya
2	2	Aarav	Espinoza	2019-11-11	NaN	Officer	David Ali	aarav.e
3	3	Aaron	Tapia	2023-01-20	NaN	Officer	Michael Reed	aaron
6	6	Abigail	Moran	2019-08-24	NaN	Officer	Caitlin Stokes	abagai
...
2987	2987	Zariah	Black	2019-11-13	NaN	Officer	Jenna Weaver	zaria
2991	2991	Zayne	Hunter	2019-11-11	NaN	Professor	Lisa Jordan	zayne
2996	2996	Zoey	Spence	2020-05-06	NaN	Senior Lecturer	Ian Price	zoey
2997	2997	Zoie	Logan	2019-04-19	NaN	Lecturer	Heidi Terry	zoie
2998	2998	Zoie	Mercado	2020-12-01	NaN	Lecturer	David Young	zoie.n

1467 rows × 48 columns



In [414]:

```
1 df_active.to_csv("./active_staff_prediction.csv",index=False)
```

In [415]:

```
1 # Select features from the whole dataframe
2 X_all = df[features]
3
4 # Predict termination class on entire data
5 pred_all = pipeline.predict(X_all)
6
7 # Predict probability for termination (class 1)
8 prob_all = pipeline.predict_proba(X_all)[:, 1]
9
10 # Add predictions and probabilities back to the dataframe
11 df['Predicted_Terminated'] = pred_all
12 df['Probability_Terminated'] = prob_all
13
14 # Evaluate performance on full dataset
15 from sklearn.metrics import classification_report, accuracy_score
16
17 print(classification_report(df['TerminatedStaff'], df['Predicted_Termin
18
19
20
```

	precision	recall	f1-score	support
0	0.96	0.97	0.97	1467
1	0.97	0.96	0.97	1533
accuracy			0.97	3000
macro avg	0.97	0.97	0.97	3000
weighted avg	0.97	0.97	0.97	3000

In [416]:

```
1 df
```

Out[416]:

Unnamed: 0		FirstName	LastName	StartDate	ExitDate	PositionTitle	Supervisor	
0	0	Aaden	Mercer	2023-07-26	NaN	Senior Research Fellow	Victoria Jacobs	aaden
1	1	Aaliyah	Watts	2020-05-09	NaN	Senior Lecturer	Jared Whitehead	aaliya
2	2	Aarav	Espinoza	2019-11-11	NaN	Officer	David Ali	aarav.e
3	3	Aaron	Tapia	2023-01-20	NaN	Officer	Michael Reed	aaron
4	4	Aaron	Weber	2020-04-06	2023-07-24	Lecturer	Melanie Garcia	aaron
...
2995	2995	Zoey	Page	2019-05-06	2022-04-09	Professor	William Newman	zoey
2996	2996	Zoey	Spence	2020-05-06	NaN	Senior Lecturer	Ian Price	zoey
2997	2997	Zoie	Logan	2019-04-19	NaN	Lecturer	Heidi Terry	zoie
2998	2998	Zoie	Mercado	2020-12-01	NaN	Lecturer	David Young	zoie.n
2999	2999	Zoie	Rowland	2022-02-08	2022-07-22	Senior Lecturer	Mrs. Rachel Henry	zoie.

3000 rows × 48 columns



In [417]:

```
1 print(f"Accuracy on full data: {accuracy_score(df['TerminatedStaff'], d
2
```

Accuracy on full data: 0.9680

In [418]:

```
1 df.to_csv("./Predict_Terminations.csv",index=False)
```

In [419]:

```
1 print(df[ ['TerminatedStaff', 'Predicted_Terminated', 'Probability_Term
```

	TerminatedStaff	Predicted_Terminated	Probability_Terminated
0	0	0	0.062233
1	0	0	0.000448
2	0	0	0.000091
3	0	0	0.044958
4	1	1	0.851542

In [420]:

```
1 # 1. Get feature names after preprocessing
2 # For OneHotEncoder, get categories and build feature names
3 ohe = pipeline.named_steps['preprocessor'].named_transformers_['cat']
4 ohe_feature_names = ohe.get_feature_names_out(categorical_features)
5
6 # Combine with numerical feature names (passed through)
7 feature_names = list(ohe_feature_names) + numerical_features
8
9 # 2. Get the trained XGBClassifier model
10 model = pipeline.named_steps['classifier']
11
12 # 3. Get feature importance scores (gain or weight or cover)
13 importance_dict = model.get_booster().get_score(importance_type='gain')
14
15 # 4. Map feature names to importance scores (note: importance keys might
16 importance_df = pd.DataFrame({
17     'Feature': feature_names,
18     'Importance': [importance_dict.get(f'f{i}', 0) for i in range(len(f
19 })
20
21 # Sort by importance descending
22 importance_df = importance_df.sort_values(by='Importance', ascending=False)
23
24 print(importance_df)
25
```

	Feature	Importance
0	EmploymentDuration_months	7.525079
1	StaffType_Academic	1.220766
2	Age	0.980411
3	Engagement Score	0.978649
4	EmployeeType_Part-Time	0.970561
5	Total_Survey_Score	0.970504
6	Faculty_Faculty of Arts	0.954568
7	Satisfaction Score	0.935931
8	Faculty_Faculty of Engineering	0.924429
9	Current Employee Rating	0.922855
10	Work-Life Balance Score	0.909402
11	PositionTitle_LE	0.900780
12	GenderCode_Female	0.894735
13	JobFunction_LE	0.881341
14	Faculty_Faculty of Medicine	0.876819
15	EmployeeType_Full-Time	0.799365
16	EmployeeType_Contract	0.794673
17	TenureGroup_	0.613771
18	GenderCode_Male	0.000000
19	StaffType_Professional	0.000000

```

In [421]: 1 import matplotlib.pyplot as plt
          2 import seaborn as sns
          3
          4 # Assuming importance_df is created as before
          5 plt.figure(figsize=(12, 8))
          6 sns.barplot(data=importance_df.head(20), x='Importance', y='Feature', p
          7 plt.title('Top 20 Feature Importances from XGBoost')
          8 plt.xlabel('Importance (Gain)')
          9 plt.ylabel('Feature')
         10 plt.tight_layout()
         11 plt.show()
         12

```

